

ElasticSearch:

下载 ElasticSearch 安装包:

Linux: https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-7.0.0-linux-x86_64.tar.gz

Windows: https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-7.0.0-windows-x86_64.zip

解压: `tar -xzf elasticsearch-7.0.0-linux-x86_64.tar.gz`

编辑 ElasticSearch 配置文件 `elasticsearch-7.0.0/config/elasticsearch.yml`, 修改以下参数内容 (示例):

```
cluster.name: ck-elasticsearch
node.name: node-01
network.host: 10.0.1.4
discovery.seed_hosts: ["127.0.0.1"]
cluster.initial_master_nodes: ["node-01"]
```

其中 `network.host` 的 ip 地址为 VM 的 `ifconfig` 得到的 `eth0` 地址 (azure VM 对外访问的 `nic` 内网地址)

切换到 `root`, 修改系统 `limit` 参数:

运行: `vi /etc/sysctl.conf`, 添加:

```
vm.max_map_count=655350
```

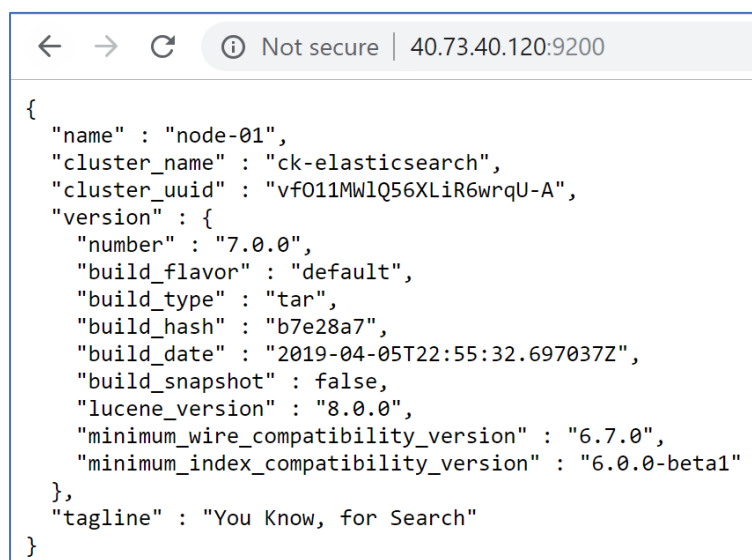
保存后运行 `sysctl -p`

运行: `vi /etc/security/limits.conf`, 加入:

```
* soft nfile 655350
* hard nfile 655350
```

重新打开一个 `ssh` 终端, 到 `elasticsearch-7.0.0/bin` 目录下, 运行: `./elasticsearch -d`

到 `azure portal` 上, 放开 ElasticSearch 所在虚机的 9200 端口的入站允许
使用浏览器或 `curl` 访问 `<VM_ip>:9200`, 应看到如下内容:



```
{
  "name" : "node-01",
  "cluster_name" : "ck-elasticsearch",
  "cluster_uuid" : "vf011MWlQ56XLiR6wrqU-A",
  "version" : {
    "number" : "7.0.0",
    "build_flavor" : "default",
    "build_type" : "tar",
    "build_hash" : "b7e28a7",
    "build_date" : "2019-04-05T22:55:32.697037Z",
    "build_snapshot" : false,
    "lucene_version" : "8.0.0",
    "minimum_wire_compatibility_version" : "6.7.0",
    "minimum_index_compatibility_version" : "6.0.0-beta1"
  },
  "tagline" : "You Know, for Search"
}
```

访问接口 http://40.73.40.120:9200/_cat/nodes?v，看到如下内容则 ElasticSearch 安装成功：

Not secure 40.73.40.120:9200/_cat/nodes?v										
ip	heap.percent	ram.percent	cpu	load_1m	load_5m	load_15m	node.role	master	name	
172.16.0.4	18	45	1	0.00	0.01	0.05	mdi	*	node-01	

Kibana:

下载 Kibana 安装包：

Linux: https://artifacts.elastic.co/downloads/kibana/kibana-7.0.0-linux-x86_64.tar.gz

Windows: https://artifacts.elastic.co/downloads/kibana/kibana-7.0.0-windows-x86_64.zip

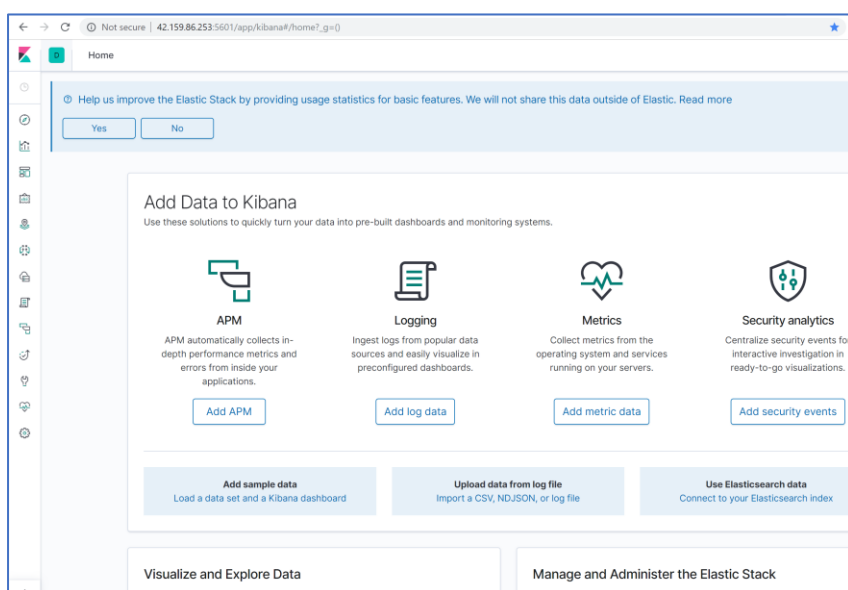
解压: `tar -xzf kibana-7.0.0-linux-x86_64.tar.gz`

编辑 ElasticSearch 配置文件 `kibana-7.0.0-linux-x86_64/config/kibana.yml`，修改以下参数内容（示例）：

```
server.port: 5601
server.host: "10.0.1.4"
elasticsearch.hosts: ["http://40.73.40.120:9200"]
kibana.index: ".kibana"
server.ssl.enabled: false
```

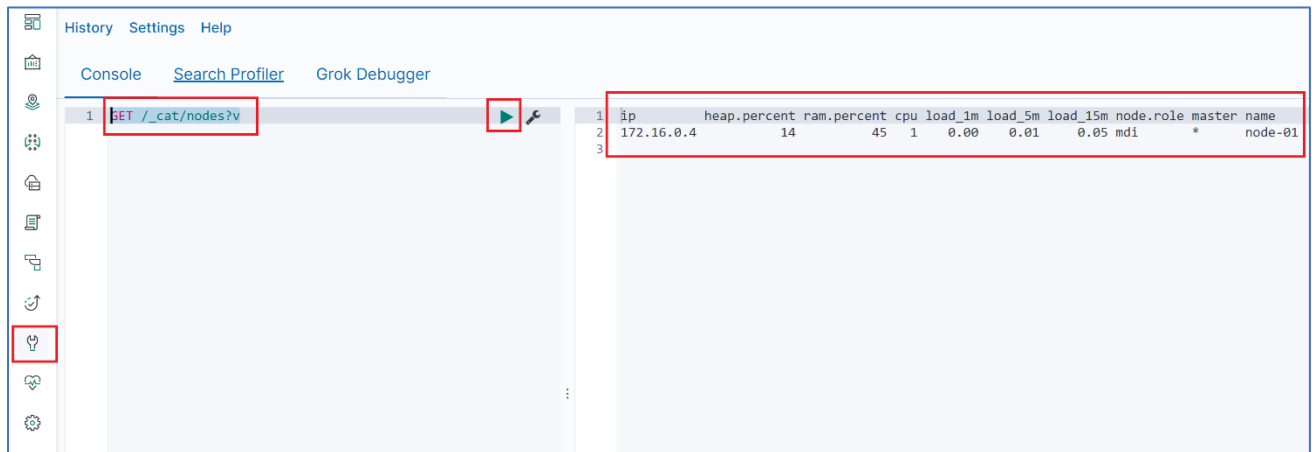
其中 `server.host` 的 ip 地址为 VM 的 `ifconfig` 得到的 `eth0` 地址（azure VM 对外访问的 `nic` 内网地址）
`elasticsearch.hosts` 为 kibana 可以访问到的 ElasticSearch 服务的地址和 9200 端口

到 azure portal 上，放开 Kibana 所在虚机的 5601 端口的入站允许
到 `kibana-7.0.0-linux-x86_64/bin` 目录下，运行 `nohup ./kibana &`
使用浏览器或 `curl` 访问 `<VM_ip>:5601`，看到如下内容则 Kibana 安装成功：



点击左侧工具栏的“Dev Tools”，执行 API “GET /_cat/nodes?v”，看到如下信息和以上 ElasticSearch 安装后执

行 API 的内容一致，则 Kibana 和 ElasticSearch 集成完毕：



Fluent-bit

执行命令创建 dev namespace （namespace 按实际需要命名）

```
kubectl create namespace dev
```

Git clone <https://github.com/fluent/fluent-bit-kubernetes-logging> 到本地

进入 `fluent-bit-kubernetes-logging\` 目录，修改以下文件中的 namespace 值为 dev：

`fluent-bit-role-binding.yaml`

`fluent-bit-service-account.yaml`

执行：

```
kubectl apply -f fluent-bit-service-account.yaml
```

```
kubectl apply -f fluent-bit-role.yaml
```

```
kubectl apply -f fluent-bit-role-binding.yaml
```

进入 `fluent-bit-kubernetes-logging\output\elasticsearch\` 目录，根据实际情况修改文件 `fluent-bit-configmap.yaml` 中以下 highlight 部分：

```
filter-kubernetes.conf: |
  [FILTER]
    Name          kubernetes
    Match         kube.*
    Kube_URL      https://ptfm-aks-u-ptfm-rsgp-uat-071a5d-3d6cbcd2.hcp.chinaeast2.cx.prod.service.azk8s.cn:443
    Kube_CA_File  /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
    Kube-Token_File /var/run/secrets/kubernetes.io/serviceaccount/token
    Merge_Log     On
    K8S-Logging.Parser On
    K8S-Logging.Exclude Off

output-elasticsearch.conf: |
```

[OUTPUT]

```
Name          es
Match         *
Host          42.159.86.253
Port          9200
Logstash_Format On
Logstash_Prefix aks-fluentbit
Replace_Dots   On
Retry_Limit    False
```

进入 fluent-bit-kubernetes-logging\output\elasticsearch\ 目录，根据实际情况修改文件 fluent-bit-ds.yaml 中以下 highlight 部分：

```
spec:
  containers:
  - name: fluent-bit
    image: fluent/fluent-bit:1.0.6
    imagePullPolicy: Always
    ports:
    - containerPort: 2020
    env:
    - name: FLUENT_ELASTICSEARCH_HOST
      value: "10.106.196.20"
    - name: FLUENT_ELASTICSEARCH_PORT
      value: "9200"
    volumeMounts:
    - name: varlog
      mountPath: /var/log
    - name: varlibdockercontainers
      mountPath: /var/lib/docker/containers
      readOnly: true
    - name: fluent-bit-config
      mountPath: /fluent-bit/etc/
    terminationGracePeriodSeconds: 10
```

执行命令：

```
kubectl apply -f fluent-bit-configmap.yaml
```

```
kubectl apply -f fluent-bit-ds.yaml
```

使用以下命令检查 fluent-bit 正常运行：

```
F:\AKS-env-setup>kubectl get daemonset -n dev
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
fluent-bit	3	3	3	3	3	<none>	52m

稍等片刻，到浏览器访问 http://10.106.196.20:9200/_cat/indices，得到以下输出：

10.106.196.20:9200/_cat/indices										
green	open	.kibana_task_manager	Ucj994TtQfqRS_YGjXm22g	1	0	2	0	45.5kb	45.5kb	
yellow	open	aks-fluentbit-2019.05.05	NKvzzjHDT0OvIqAmLuhZUA	1	1	21673	0	7mb	7mb	
green	open	.kibana_1	O5bgIhzdQwCbug6d3_BL9g	1	0	6	0	35.3kb	35.3kb	

其中 aks-fluentbit-2019.05.05 是 fluentbit 收集 AKS 节点上的容器日志，在 ElasticSearch 中创建索引并上传日志内容。

至此 fluentbit 配置完毕。

Filebeat

注意 filebeat 和 fluent-bit 都是用来收集日志，并发送到 ElasticSearch，因此运行两者之一就可以。

参考 yaml 文件：

<https://github.com/kylercai/kylercRepo/blob/master/aks/efk-logging/filebeat-kubernetes.yaml>

此 yaml 文件定义以 daemonset 的方式运行 filebeat, 同时创建运行 filebeat daemonset 所需的 service account、cluster role、cluster role binding、configmap 等。

参考以下 highlight 部分对 yaml 文件内容进行定制，其中 multiline.*内容定义 filebeat 对多行日志进行合并显示。

```
---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: filebeat
  namespace: kube-system
  labels:
    k8s-app: filebeat
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRole
metadata:
  name: filebeat
  labels:
    k8s-app: filebeat
rules:
- apiGroups: ["" ] # "" indicates the core API group
  resources:
  - namespaces
```

```

- pods
verbs:
- get
- watch
- list
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRoleBinding
metadata:
  name: filebeat
subjects:
- kind: ServiceAccount
  name: filebeat
  namespace: kube-system
roleRef:
  kind: ClusterRole
  name: filebeat
  apiGroup: rbac.authorization.k8s.io
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: filebeat-config
  namespace: kube-system
  labels:
    k8s-app: filebeat
data:
  filebeat.yml: |-
    filebeat.inputs:
      - type: log
      enabled: true
      paths:
        - /var/log/containers/*.log

    filebeat.config:
      inputs:
        # Mounted `filebeat-inputs` configmap:
        path: ${path.config}/inputs.d/*.yaml
        # Reload inputs configs as they change:
        reload.enabled: false
      modules:
        path: ${path.config}/modules.d/*.yaml
        # Reload module configs as they change:
        reload.enabled: false

```

```

# To enable hints based autodiscover, remove `filebeat.config.inputs` configuration and uncomment this:
#filebeat.autodiscover:
# providers:
#   - type: kubernetes
#     hints.enabled: true

processors:
  - add_cloud_metadata:

output.elasticsearch:
  hosts: ["42.159.86.253:9200"]
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: filebeat-inputs
  namespace: kube-system
  labels:
    k8s-app: filebeat
data:
  kubernetes.yml: |-
    - type: docker
      containers.ids:
        - "*"
    multiline.pattern: '^[[:space:]]|^Caused by:'
    multiline.negate: false
    multiline.match: after
  processors:
    - add_kubernetes_metadata:
        in_cluster: true
---
apiVersion: extensions/v1beta1
kind: DaemonSet
metadata:
  name: filebeat
  namespace: kube-system
  labels:
    k8s-app: filebeat
spec:
  template:
    metadata:
      labels:
        k8s-app: filebeat

```

```
spec:
  serviceAccountName: filebeat
  terminationGracePeriodSeconds: 30
  containers:
  - name: filebeat
    image: docker.elastic.co/beats/filebeat:7.1.0
    args: [
      "-c", "/etc/filebeat.yml",
      "-e",
    ]
    env:
    - name: ELASTICSEARCH_HOST
      value: "42.159.86.253"
    - name: ELASTICSEARCH_PORT
      value: "9200"
    securityContext:
      runAsUser: 0
      # If using Red Hat OpenShift uncomment this:
      #privileged: true
    resources:
      limits:
        memory: 200Mi
      requests:
        cpu: 100m
        memory: 100Mi
    volumeMounts:
    - name: config
      mountPath: /etc/filebeat.yml
      readOnly: true
      subPath: filebeat.yml
    - name: inputs
      mountPath: /usr/share/filebeat/inputs.d
      readOnly: true
    - name: data
      mountPath: /usr/share/filebeat/data
    - name: varlibdockercontainers
      mountPath: /var/lib/docker/containers
      readOnly: true
  volumes:
  - name: config
    configMap:
      defaultMode: 0600
      name: filebeat-config
  - name: varlibdockercontainers
```



```

    hostPath:
      path: /var/lib/docker/containers
  - name: inputs
    configMap:
      defaultMode: 0600
      name: filebeat-inputs
  # data folder stores a registry of read status for all files, so we don't send everything again on a
  Filebeat pod restart
  - name: data
    hostPath:
      path: /var/lib/filebeat-data
      type: DirectoryOrCreate
---

```

其中 multiline 设置:

```

multiline.pattern: '^[[:space:]]|^Caused by:'
multiline.negate: false
multiline.match: after

```

含义为:

符合 pattern 的行, 即以空格和"Caused by:"开头的行, 将被合并到上一行之后。

多行日志合并设置应根据自身的日志合并需要调整。multiline 还有以下 option 可供调整:

- multiline.flush_pattern
指定正则表达式去匹配指定的行作为 multiline-message 的结束, 刷新的内存, 开始匹配新的多行
- multiline.max_lines
指定合并最大行数
- multiline.timeout
设定一个超时时间, 在时间结束后, 即使没有匹配到新 pattern 来启动新事件, Filebeat 也会发送多行事件。默认值是 5 秒

执行命令:

```

λ kubectl apply -f "filebeat-kubernetes.yaml"
serviceaccount/filebeat created
clusterrole.rbac.authorization.k8s.io/filebeat created
clusterrolebinding.rbac.authorization.k8s.io/filebeat created
configmap/filebeat-config created
configmap/filebeat-inputs created
daemonset.extensions/filebeat created

```

使用以下命令检查 filebeat 正常运行：

```
λ kubectl get daemonset -n kube-system
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
azure-cni-networkmonitor	3	3	3	3	3	beta.kubernetes.io/os=linux	29d
azure-ip-masq-agent	3	3	3	3	3	beta.kubernetes.io/os=linux	29d
filebeat	3	3	3	3	3	<none>	12h
kube-proxy	3	3	3	3	3	beta.kubernetes.io/os=linux	29d
kube-svc-redirect	3	3	3	3	3	beta.kubernetes.io/os=linux	29d
omsagent	3	3	3	3	3	beta.kubernetes.io/os=linux	29d

稍等片刻，到浏览器访问 http://42.159.86.253:9200/_cat/indices，得到以下输出（highlight filebeat 在 ElasticSearch 中创建的 index，并用于上传日志）：

yellow	open	aks-fluentbit-2019.05.20	aGZkItrhRDkDbDXf1Z-JdA	1	1	914916	0	173.5mb	173.5mb
yellow	open	aks-fluentbit-2019.05.21	00giDzwxRBadqIf8kmmRng	1	1	415048	0	80.1mb	80.1mb
yellow	open	aks-fluentbit-2019.05.13	bKH-XG6NSauUIJuZmvMO7w	1	1	915242	0	174.7mb	174.7mb
yellow	open	aks-fluentbit-2019.05.08	ya2HY8bMRU2W8R-ykh1NOQ	1	1	915047	0	173.4mb	173.4mb
yellow	open	aks-fluentbit-2019.05.18	yI7HrqdKQ4iuypfVFiINKA	1	1	915854	0	174.4mb	174.4mb
yellow	open	aks-fluentbit-2019.05.06	jvE20VkCTUmMGe1LNrQQAA	1	1	914914	0	173.9mb	173.9mb
yellow	open	aks-fluentbit-2019.05.10	08VtEU5bTN-4AV0ob63xuQ	1	1	914462	0	172.4mb	172.4mb
yellow	open	aks-fluentbit-2019.05.19	o-XPtmK0STSEi8AJ8SrBgg	1	1	914903	0	174.1mb	174.1mb
yellow	open	aks-fluentbit-2019.05.14	BMNL74mtS00hYLYzlskGFA	1	1	914347	0	172.9mb	172.9mb
green	open	.kibana_1	wE4S--JuRiyy1EiqmOZZbw	1	0	6	0	91.1kb	91.1kb
yellow	open	aks-fluentbit-2019.05.11	kwMNNdpXTKqnBro-UlSPRQ	1	1	914825	0	173.8mb	173.8mb
yellow	open	aks-fluentbit-2019.05.15	LEubQZRbQGEXTiwfYA_MYw	1	1	914102	0	173.9mb	173.9mb
yellow	open	aks-fluentbit-2019.05.05	bpv6MaZrTHmQiMgSF1FzvA	1	1	370304	0	70.1mb	70.1mb
yellow	open	aks-fluentbit-2019.05.12	CZNCR3C7T4mwE1N32xuMFQ	1	1	914356	0	173.5mb	173.5mb
yellow	open	aks-fluentbit-2019.05.07	PvNq0DFbTEC18MSU5_4Wyg	1	1	916966	0	175.1mb	175.1mb
green	open	.kibana_task_manager	WYZPf_WIQ7evnSRg6ZUCWg	1	0	2	0	21.4kb	21.4kb
yellow	open	aks-fluentbit-2019.05.16	IXPUA-DARE2DerxjNvFizg	1	1	916239	0	175.2mb	175.2mb
yellow	open	aks-fluentbit-2019.05.09	NGeKiWwCtpm7XiVi0w22mw	1	1	914482	0	173.4mb	173.4mb
yellow	open	aks-fluentbit-2019.05.17	7XBiuGmMScSycmrc-nh7zA	1	1	914555	0	173.3mb	173.3mb
yellow	open	filebeat-7.1.0-2019.05.21-000001	OURsv5NsQ7aygBjbycyn1A	1	1	1245225	0	390.8mb	390.8mb

至此 filebeat 配置完毕。

配置 Kibana 查询日志

进入 Kibana 首页面板，选择左侧工具栏最下方 Management（齿轮图标）

点击 ElasticSearch -> Index Management，应能看到 ElasticSearch 中已有的索引

点击 Kibana -> Index Patterns -> Create index pattern，创建索引模式，例如：

Create index pattern

★ aks-fluentbit*

Create index pattern

Kibana uses index patterns to retrieve data from Elasticsearch indices for things like

Step 1 of 2: Define index pattern

Index pattern

aks-fluent*

You can use a * as a wildcard in your index pattern.
You can't use spaces or the characters \, /, ?, ", <, >, |.

✓ **Success!** Your index pattern matches **1 index**.

aks-fluentbit-2019.05.05

Rows per page: 10 ▾

如果日志收集端选择 filebeat，则在此处的 index pattern 可以选择“filebeat-*”。

点击 Next step，进行配置后点击 Create index pattern，例如：

Create index pattern

Kibana uses index patterns to retrieve data from Elasticsearch indices for things like visualizations.

☐ Include system indices

Step 2 of 2: Configure settings

You've defined **aks-fluent*** as your index pattern. Now you can specify some settings before we create it.

Time Filter field name

Refresh

@timestamp ▾

@timestamp

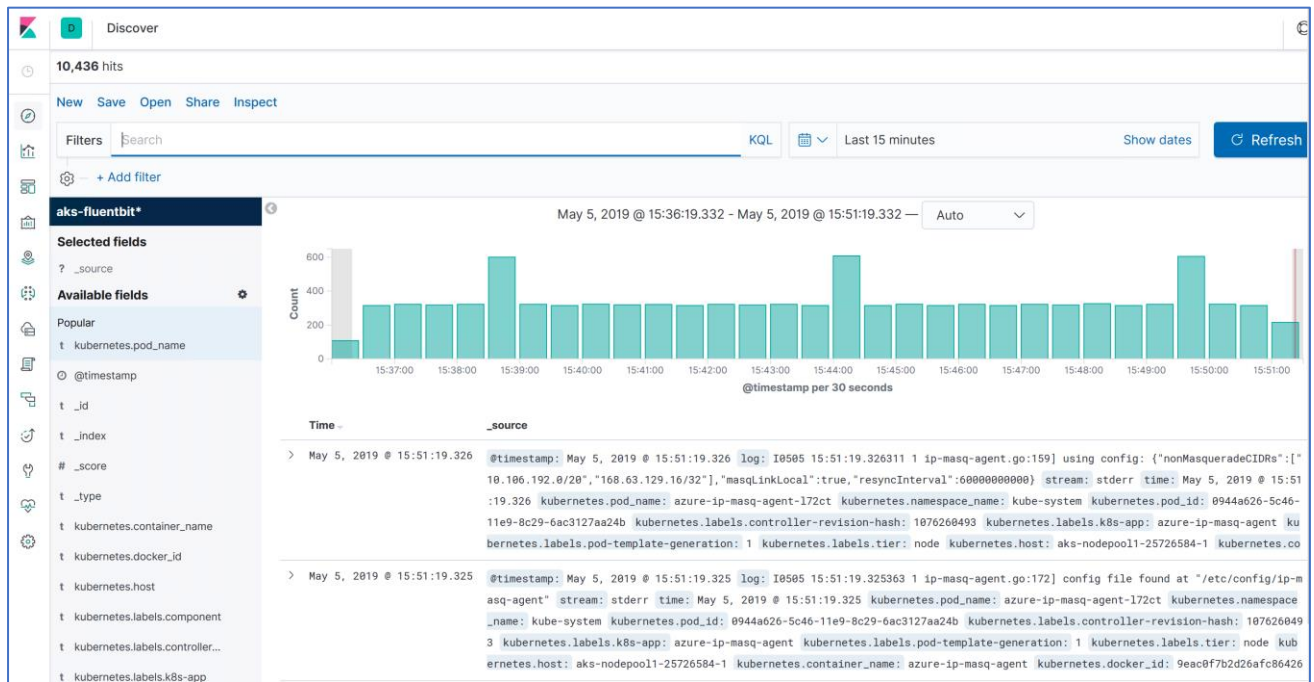
time

I don't want to use the Time Filter

< Back

Create index pattern

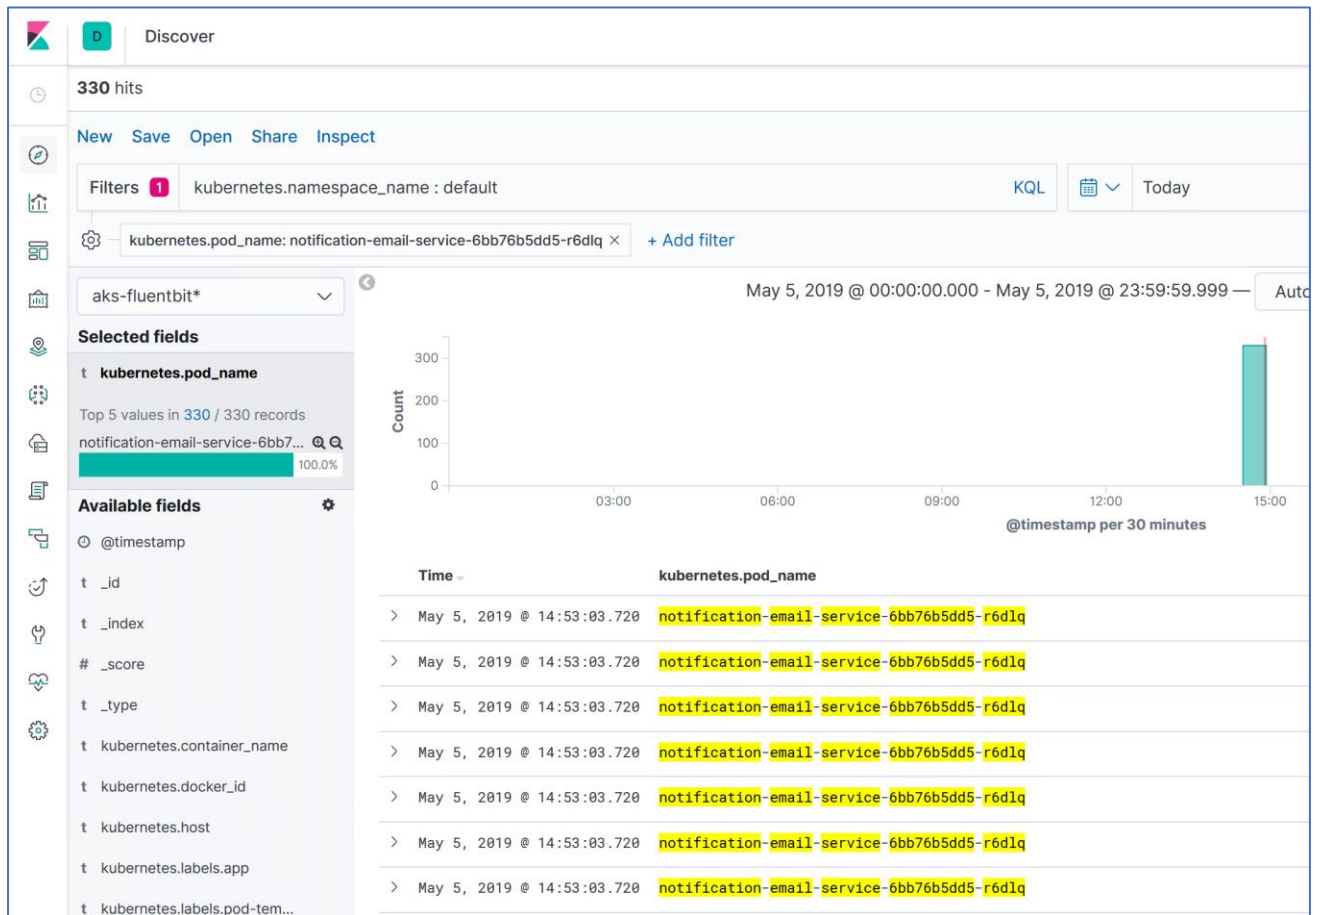
之后，到工具栏左侧选择 Discover（最上方图标），选择对应的 index pattern，Kibana 将自动获取相应的日志记录，例如：



根据查询需要，设置过滤条件，缩小日志查询范围，可以使用的方式有：

- 在 Filters 中输入查询条件，例如 `kubernetes.namespace_name : default` 为查询 kubernetes 中 default namespace 中的内容（请查阅 KQL 查询语法）
- 在 Available fields 中移动鼠标到相应 fields，点击 add，相应的 field 会被添加到 Selected fields，点击此 field 会显示所有的 value，选择相应的取值进行过滤

例如：



过滤条件选择了：
default namespace
Kubernetes pod 名字为 notification-email-service-xxxxx

如果日志收集端选择 filebeat，并且配置了 multiline.pattern 等多行日志合并选项，则可在 Kibana 中看到相关多行日志合并显示的效果：

标准输出 stdout：

```
t input.type                docker
t kubernetes.container.name javademo
t kubernetes.labels.app     javademo
t kubernetes.labels.pod-template-hash 3527011171
t kubernetes.namespace      default
t kubernetes.node.name       aks-agentpool-32361114-2
t kubernetes.pod.name        javademo-796c4555c5-xmptr
t kubernetes.pod.uid         11358850-7bdf-11e9-9ffe-b226bbbe90d4
t kubernetes.replicaset.name javademo-796c4555c5
t log.file.path              /var/lib/docker/containers/ee0f6c822333a2e5e76d51119a7d5b01d2e0eb0006b7ddb0361efbc8c8a813e4/ee0f6c822333a2e5e76d51119
1d2e0eb0006b7ddb0361efbc8c8a813e4-json.log
t log.flags                  multiline
# log.offset                 24,601
t message                    This log simulates the multiple line log info collected&combined by filebeats..
                             line 1 begins with space..
                             line 2 begins with space..
                             line 3 begins with space..
                             at line 4 begins with space followed by word at..
                             ... line 5 begins with space followed by ...
                             Caused by: line 6 begins with words Caused by:..
t stream                     stdout
○ suricata.eve.timestamp     May 22, 2019 @ 00:33:06.464
```

```
▼ May 22, 2019 @ 00:33:06.464 kubernetes.namespace: default @timestamp: May 22, 2019 @ 00:33:06.464 ecs.version: 1.0.0 agent.version: 7.1.0 agent.type: filebeat
agent.ephemeral_id: b4c3ebda-b01a-4030-9396-99126732a151 agent.hostname: filebeat-9d5b4 agent.id: bfdbbe87-9651-4c52-9011-f0a170abdf7b
message: java.io.IOException: This log simulates an exception thrown.. at
com.microsoft.azure.webapp.HelloWorldServlet.doGet(HelloWorldServlet.java:51) at javax.servlet.http.HttpServlet.service(HttpServlet.java:6
javax.servlet.http.HttpServlet.service(HttpServlet.java:742) at
```

Expanded document

[View surrounding documents](#) [View single document](#)

标准错误 stderr:

```
t kubernetes.replicaset.name javademo-796c4555c5
t log.file.path              /var/lib/docker/containers/ee0f6c822333a2e5e76d51119a7d5b01d2e0eb0006b7ddb0361efbc8c8a813e4/ee0f6c822333a2e5e76d51119
1d2e0eb0006b7ddb0361efbc8c8a813e4-json.log
t log.flags                  multiline
# log.offset                 25,410
t message                    java.io.IOException: This log simulates an exception thrown..
                             at com.microsoft.azure.webapp.HelloWorldServlet.doGet(HelloWorldServlet.java:51)
                             at javax.servlet.http.HttpServlet.service(HttpServlet.java:635)
                             at javax.servlet.http.HttpServlet.service(HttpServlet.java:742)
                             at org.apache.catalina.core.ApplicationFilterChain.internalDoFilter(ApplicationFilterChain.java:231)
                             at org.apache.catalina.core.ApplicationFilterChain.doFilter(ApplicationFilterChain.java:166)
                             at org.apache.tomcat.websocket.server.WsFilter.doFilter(WsFilter.java:52)
                             at org.apache.catalina.core.ApplicationFilterChain.internalDoFilter(ApplicationFilterChain.java:193)
                             at org.apache.catalina.core.ApplicationFilterChain.doFilter(ApplicationFilterChain.java:166)
                             at org.apache.catalina.core.StandardWrapperValve.invoke(StandardWrapperValve.java:199)
                             at org.apache.catalina.core.StandardContextValve.invoke(StandardContextValve.java:96)
                             at org.apache.catalina.authenticator.AuthenticatorBase.invoke(AuthenticatorBase.java:493)
                             at org.apache.catalina.core.StandardHostValve.invoke(StandardHostValve.java:137)
                             at org.apache.catalina.valves.ErrorReportValve.invoke(ErrorReportValve.java:81)
                             at org.apache.catalina.valves.AbstractAccessLogValve.invoke(AbstractAccessLogValve.java:660)
                             at org.apache.catalina.core.StandardEngineValve.invoke(StandardEngineValve.java:87)
                             at org.apache.catalina.connector.CoyoteAdapter.service(CoyoteAdapter.java:343)
                             at org.apache.coyote.http11.Http11Processor.service(Http11Processor.java:798)
                             at org.apache.coyote.AbstractProcessorLight.process(AbstractProcessorLight.java:66)
                             at org.apache.coyote.AbstractProtocol$ConnectionHandler.process(AbstractProtocol.java:808)
                             at org.apache.tomcat.util.net.NioEndpoint$SocketProcessor.doRun(NioEndpoint.java:1498)
                             at org.apache.tomcat.util.net.SocketProcessorBase.run(SocketProcessorBase.java:49)
                             at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
                             at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
                             at org.apache.tomcat.util.threads.TaskThread$WrappingRunnable.run(TaskThread.java:61)
                             at java.lang.Thread.run(Thread.java:748)
t stream                     stderr
○ suricata.eve.timestamp     May 22, 2019 @ 00:33:06.464
```

```
> May 22, 2019 @ 00:33:05.311 kubernetes.namespace: default @timestamp: May 22, 2019 @ 00:33:05.311 stream: stdout message: This log simulates the multiple line log
collected&combined by filebeats.. line 1 begins with space.. line 2 begins with space.. line 3 begins with space.. at line 4 begins with s
followed by word at.. ... line 5 begins with space followed by ... Caused by: line 6 begins with words Caused by:.. ecs.version: 1.0.0
agent.version: 7.1.0 agent.type: filebeat agent.ephemeral_id: b4c3ebda-b01a-4030-9396-99126732a151 agent.hostname: filebeat-9d5b4
agent.id: bfdbbe87-9651-4c52-9011-f0a170abdf7b log.offset: 19,693
```